

## Today's Question

How to implement parallel computing structures in C++?

### Serial Workflow



[1]

### Parallel Workflow



[1]

# Pragmatic Introduction to Intel's Threading Building Blocks

Peter Steinbach

Institute for Nuclear and Particle Physics, TU Dresden

June 20th, 2012

under CC license  BY 3.0

## A (bit of) Introduction

### Let's assume ...

- ▶ Your software **takes a too long** to complete.
- ▶ **You have identified a very CPU intense region in your code!**
- ▶ What to do? Let's go **parallel**.

### Multiple Processes

A process<sup>[2]</sup> is an instance of a computer program that is being executed. It contains the program code and its current activity. A process is said to own the following resources: image of the executable machine code, memory (stack, heap), ...

## A (bit of) Introduction

Let's assume ...

- ▶ Your software **takes a too long** to complete.
- ▶ **You have identified a very CPU intense region in your code!**
- ▶ What to do? Let's go **parallel**.

### Multiple Processes

A process<sup>[2]</sup> is an instance of a computer program that is being executed. It contains the program code and its current activity. A process is said to own the following resources: image of the executable machine code, memory (stack, heap), ...

### Multiple Threads

A thread<sup>[3]</sup> of execution is the smallest unit of processing that can be scheduled by an operating system. A thread is a lightweight process. A thread is contained inside a process. Multiple threads can exist within the same process and share resources (memory, ...).


# Parallel $\neq$ Parallel

Does the computer architecture play a role?


For Multi-Threading and Multi-Processing: No.

For your problem: Yes!

## *Shared Memory systems*



## *Shared/Distributed Memory systems*



## Let's go for Multi-Threading

I assume ...

- ▶ your serial algorithm cannot be optimised further
- ▶ problem domain can be decomposed
- ▶ speed-up through multi-threading of code is high to justify coding effort

# Let's go for Multi-Threading

I assume ...

- ▶ your serial algorithm cannot be optimised further
- ▶ problem domain can be decomposed
- ▶ speed-up through multi-threading of code is high to justify coding effort

## Always remember

Implementing concurrency (i.e. multi-threading) is quite an effort!



[4]

# What library to pick?

## Explicit Threading Libraries

Provides low level primitives (such as thread objects etc.)

- ▶ C++03 does support threads
  - ▶ [Boost.Threads](#)
  - ▶ [POSIX threads](#)
- ▶ C++11 does `std::thread`

## Implicit Threading Libraries

Provide higher level API that encapsulate thread control

- ▶ [OpenMP](#)
- ▶ [Apple's Grand Central Dispatch](#)
- ▶ [Intel Threading Building Blocks](#)

## For me as a scientist ...

by first approximation, implicit libraries are my first choice (hide threading details, provide API for most common problems).



# Intel Threading Building Blocks

## Generalities

- ▶ Open Source C++ library ([GPLv2](#))
- ▶ sponsored by Intel<sup>TM</sup>
- ▶ runs on shared memory x86 architectures
- ▶ runs with Windows, Mac and Linux (MS Visual++, icc, gcc, ... )
- ▶ current version: 4.0

## Design

- ▶ generic template library
- ▶ design similar to STL (templated algorithms interfaced with templated container types)
- ▶ developer friendly (exceptions, compiler checks type safety)
- ▶ good documentation ([tutorials](#), [Examples](#), [Code reference](#) available online)
- ▶ some low-level features: custom memory allocators, atomic operation identifiers ...

## A Quick Look Inside

### Containers

- ▶ `tbb::concurrent_vector`
- ▶ `tbb::concurrent_hash_map`
- ▶ `tbb::concurrent_queue`

### Algorithms

- ▶ `tbb::parallel_do`
- ▶ `tbb::parallel_for (*)`
- ▶ `tbb::parallel_sort`
- ▶ `tbb::parallel_reduce (*)`
- ▶ `tbb::parallel_pipeline`
- ▶ ...

## Example 1: A simple logger

### Serial Version

```
for(int i = 0; i < nIterations; i++){  
    std::cout << "iteration " << i << std::endl;  
}
```

## Example 1: A simple logger

### Serial Version

```
for(int i = 0; i < nIterations; i++){
    std::cout << "iteration_" << i << std::endl;
}
```

### Parallel

```
//...
int grainsize = nIterations/nThreads;
tbb::parallel_for(tbb::blocked_range<size_t>(0,
                                              nIterations,
                                              grainsize),
                  NumPrinter()
                  );
//...
```

## Example 2: A throwing random numbers

### Parallel Reduce

```
Sum sumWorker(data);  
tbb::parallel_reduce(tbb::blocked_range<size_t>(0,  
                                                    niterations,  
                                                    grainsize),  
                    sumWorker);
```

### Difference to Example 1

- ▶ Sum has no const operator()
- ▶ Sum has dedicated constructor with `tbb::split` macro

# Attention



Bundesarchiv, Bild 103-1094-0014-033  
Foto: Thiem, Wolfgang | 14. August 1984

[1]

Always ...

Compare Parallel and Serial Performance!  
(was/is the effort worthwhile?)

## Summary

### My Conclusion

- ▶ Multi-threading is NO SILVER BULLET
- ▶ TBB is a well composed template library
- ▶ its open source (so scientists can go for it)
- ▶ it offers high-level functionality that encapsulates low level thread management
- ▶ it offers an API to approach/solve most problems with parallel solutions
- ▶ it has a learning curve

### further reading I recommend

- ▶ [online](#) TBB documentations
- ▶ [video lecture](#) series by Clay Breshears

# References

- [1] [commons.wikimedia.org](https://commons.wikimedia.org).  
Wikicommons.
- [2] [Wikipedia: Processes](#).  
Wikipedia.
- [3] [Wikipedia: Threads](#).  
Wikipedia.
- [4] [openclipart.org](https://openclipart.org).  
OpenClipart library.



## License

This talk is published under the creative commons BY licence 3.0. For details, see [.](#)

All contained material was created by the author or obtained from creative common licensed sources (as indicated). For questions, concerns, criticism and improvements, contact: [.](#)